



B4M36SWA: Software Architectures

Deployment of traditional Java EE applications in the cloud

Jiri Pechanec
QA Engineer
May 2nd, 2017

AGENDA

- Java EE Container
 - Standard deployment
 - Clustered deployment
- Keeping the state
- Deployment in cloud
 - VMs
 - Containers
 - Scaling
 - Manual
 - Automatic

Java EE Application

- Java application running in Java EE container
- Application
 - Stateful 3-tier
 - Data Tier – standalone SQL Database
 - Logic Tier – EJBs, CDI beans
 - Presentation Tier – JSF, Servlets, JS frameworks

Java EE Application Deployment




Java EE Container

- Java runtime that provides services for Java EE applications
- Supports multi-application deployment
 - Not too used – mostly as application modules
- An application archive is placed into a deployment directory
 - EAR
 - WAR
 - JAR

Java EE Container

- Clustering
- Usually a static number of replicas on bare-metal or virtual machines
- Discovery problem
 - Static list of IPs/hostnames
 - Service registry
 - UDP Multicast

Typical clustered deployment



Typical clustered deployment

- Database should be replicated
 - Single point of failure
- Load balancer routes web client requests to an instance of application server
 - random
 - round-robin
 - load
- App server instances contains state
 - Must be replicated between instances
 - Sticky session optimization

What is state

- Typically associated with a web session
- Represents an interaction between application and the user
 - Shopping cart
 - User session
- An application server itself can keep an internal state
 - HTTP Session
 - Caches for JPA
 - Stateful Session Bean instances

Application deployment process

- Developer commits
- A final integration build
- Testing in a stage environment
- Passing the application artifacts to operations
- Outage
 - Cluster is stopped
 - A new version of the application is copied into the deployment directory
 - Start the cluster

Moving to cloud

Dynamically provided VMs

- First stage of moving to cloud
- Needs automatic cluster establishment and discovery
 - WildFly/JBoss EAP – UDP multicast
 - Load balancer must also support dynamic configuration
 - Cluster instances advertises their presence to load balancer
- Cloud provider provisions a new VM
- A VM is configured
- App server is started and automatically joins the cluster

VM Configuration

- Manually
- Automatically
 - Ansible, Puppet, ...
- Changes
 - Install the app server
 - Configure the server
 - DataSources, JDBC Driver, JCA connectors
- Deploy the application
- Time consuming and error prone process

Pre built-VM images

- Part of the application release process
- When a new version is released a pre-configured VM is built
 - Amazon AMI building
 - Oz for OpenStack
 - packer – multiplatform
- Adding a new image to cluster means starting the customized image

Container as a Service

- Cluster instances are running in Docker containers
- Similar to pre-built VM, just Docker image is built
- Allows much higher application density than VM-based solution

Platform as a Service

- The platform takes care of the whole release process
- The platform is capable of executing Maven builds
 - A location to an SCM is provided
 - A build can be triggered on-demand or automatically
- Automatic application re-deployment when a new container image is available

Java EE Application in OpenShift

Main concerns

- Automated build
- No outage deployment
- Dynamic clustering
- On-demand scaling
- Keeping the state

Automated build

- Source-2-Image (s2i) images are used to build to application
 - WildFly, JBoss EAP
 - Tomcat
- S2I executes Maven build and bakes the application into the image with Java EE container
- BuildConfig points to application sources
 - Manual trigger
 - Webhook trigger
 - Rebuild when new s2i image is available (CVEs)
- DeploymentConfig points to the image built by OpenShift
- Automatic application re-deployment when a new release is available

No outage deployment

- DeploymentConfig points to the image built by OpenShift
- Automatic application re-deployment when a new release is available
 - recreate – outage
- Rolling update
 - new nodes are started and are joining the cluster
 - old instances are removed from the cluster
 - active sessions are migrated between new and old nodes
 - Requires forward and backward compatibility!!!

Scaling

- Manual
 - `oc dc <> --replicas=`
- Automatic
 - Horizontal pod autoscaler object
 - Web UI – deployment → Actions → Add autoscaler
 - Minimum and maximum of pods
 - CPU threshold
 - Prerequisites
 - OpenShift Metrics deployed in cluster
 - Resource limits for CPU set

Keeping the state

- Database
 - no-brainer
 - requires persistent volume for database datafiles
 - MySQL, PostgreSQL, MongoDB
- Replicating the state between WildFly/JBoss EAP instances
 - Application must be <distributed/>
 - Internally uses Infinispan project – a shared cache

Database deployment

- Deploy database exposed as a service <name>-<databasetype>
- Deploy a JBoss EAP with env vars
 - DB_SERVICE_PREFIX_MAPPING: <prefix>= <name>-<databasetype>, ...
 - <prefix>_JNDI
 - <prefix>_USERNAME
 - <prefix>_PASSWORD
 - <prefix>_DATABASE
 - TX_DATABASE_PREFIX_MAPPING
 - transaction manager log storage
- Multiple databases supported

Transaction Manager concerns

- Transaction Manager controls 2PC transactions
- Requires a transaction log to bookkeep transactions in progress
- Problem with scaling in case of using file system
 - Nodes come and go
- File-based approach with database bookkeeping

KUBE_PING

- Service discovery
 - KUBE_PING protocol for JGroups networking stack
- Requires *view* privilege for service account or project
 - `oc policy add-role-to-user view system:serviceaccount:$(oc project -q):<account name|default> -n $(oc project -q)`
- Env vars for pod must be set
 - `OPENSIFT_KUBE_PING_NAMESPACE` – name of the project
 - `OPENSIFT_KUBE_PING_LABELS` – labels to identify pods to be merged in a cluster
- Container port 8888 must be exposed

KUBE_PING

- Service registry based discovery
- OpenShift provides a list of pods based on label filter
- Each pod exposes a simple HTTP server on port 8888 for cluster forming communication

Supplementary services

- Remote cache
 - A standalone running instance of Infinispan (JBoss Data Grid)
 - REST
 - HotRod
 - memcached
 - Application stores its state here instead of a database
 - Hibernate OGM
- Messaging
 - Java Message Service

Java Message Service

- Asynchronous communication
- Decouples source and target
- Two models
 - queues – point-to-point
 - topics – publish/subscribe

HornetQ

- A component of JBoss EAP
- Automatically clustered
- Requires env vars
 - MQ_CLUSTER_PASSWORD
 - MQ_QUEUES
 - MQ_TOPICS
- Can deliver messages between pods in the cluster

JBoss A-MQ

- A standalone message broker, multiprotocol
 - Openwire, AMQP, MQTT, STOMP
- Based on Apache ActiveMQ
- Static scaling only
- JBoss EAP can automatically replace an internal HornetQ with a link to a standalone JBoss A-MQ

JBoss A-MQ deployment

- Deploy a JBoss A-MQ instance, exposed as a service <name>-amq
- Deploy a JBoss EAP with env vars
 - MQ_SERVICE_PREFIX_MAPPING: <prefix>= <name>-<databasetype>, ...
 - <prefix>_JNDI
 - <prefix>_USERNAME
 - <prefix>_PASSWORD
 - <prefix>_PROTOCOL
 - <prefix>_QUEUES
 - <prefix>_TOPICS
- Multiple brokers supported

Summary

Summary

- Elasticity of cloud bot simplifies and complicates application deployment
- Key differences between traditional and cloud deployment
 - Deployment pipeline
 - Elastic scaling
 - Application state management



redhat.

THANK YOU



plus.google.com/+RedHat



facebook.com/redhatinc



linkedin.com/company/red-hat



twitter.com/RedHatNews



youtube.com/user/RedHatVideos